

```
pip install dgl -f https://data.dgl.ai/wheels/torch-2.3/cu121/repo.html
```

```
#os.environ['CUDA_LAUNCH_BLOCKING'] = '1'
```

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.nn.functional as F
import dgl
import os
from dgl.nn import GraphConv
import numpy as np
import torch.optim as optim
from torch.optim import Adam
from sklearn.model_selection import train_test_split
from torch.utils.data import Dataset
from dgl.data.loading import GraphDataLoader
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import make_scorer, mean_squared_error
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report, confusion_matrix, roc_curve, auc, roc_auc_score
from joblib import dump, load
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.linear_model import TweedieRegressor
from sklearn.neural_network import MLPRegressor
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.preprocessing import label_binarize
```

```
import warnings
warnings.filterwarnings("ignore")
```

```
print(torch.__version__)
```

```
🔄 2.3.1+cu121
```

```
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print('Using device:', device)
```

```
🔄 Using device: cuda
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
dirve_path = '/content/drive/MyDrive/Project/'
train = pd.read_csv(dirve_path + 'security_train.csv')
```

train

	file_id	label	api	tid	index
0	1	5	LdrLoadDll	2488	0
1	1	5	LdrGetProcedureAddress	2488	1
2	1	5	LdrGetProcedureAddress	2488	2
3	1	5	LdrGetProcedureAddress	2488	3
4	1	5	LdrGetProcedureAddress	2488	4
...
89806688	13887	2	NtClose	2336	618
89806689	13887	2	NtClose	2336	619
89806690	13887	2	NtClose	2336	620
89806691	13887	2	NtClose	2336	621
89806692	13887	2	NtTerminateProcess	2336	622

89806692 rows × 5 columns

```
print(train.isnull().sum())
```

```
train['file_id'] = train['file_id'].astype(np.int16)
train['label'] = train['label'].astype(np.int16)
train['api'] = train['api'].astype(str)
train['tid'] = train['tid'].astype(np.int16)
train['index'] = train['index'].astype(np.int16)
```

DataLoader

```
api_to_idx = {api: idx for idx, api in enumerate(train['api'].unique())}
```

```
train['api_idx'] = train['api'].map(api_to_idx) # Map API names to indexes
train.sort_values(by=['file_id', 'index'], inplace=True) # Ensure that the data within each file_id is sorted by index
file_ids = train['file_id'].unique()
```

```
train_file_ids, val_file_ids = train_test_split(file_ids, test_size=0.2, random_state=42)
train_data = train[train['file_id'].isin(train_file_ids)]
val_data = train[train['file_id'].isin(val_file_ids)]
```

```

train_file_ids, temp_file_ids = train_test_split(file_ids, test_size=0.3, random_state=42)

val_file_ids, test_file_ids = train_test_split(temp_file_ids, test_size=0.5, random_state=42)

train_data = train[train['file_id'].isin(train_file_ids)]
val_data = train[train['file_id'].isin(val_file_ids)]
test_data = train[train['file_id'].isin(test_file_ids)]

```

```

class CustomGraphDataset(Dataset):
    def __init__(self, dataframe, api_to_idx, device):
        self.dataframe = dataframe
        self.api_to_idx = api_to_idx
        self.file_ids = dataframe['file_id'].unique()
        self.device = device

    def __len__(self):
        return len(self.file_ids)

    def __getitem__(self, idx):
        file_id = self.file_ids[idx]
        sub_df = self.dataframe[self.dataframe['file_id'] == file_id].reset_index(drop=True)

        g = dgl.graph([], []), device=self.device) # Create an empty diagram
        g.add_nodes(len(sub_df))

        apis_indices = [self.api_to_idx[api] for api in sub_df['api']]
        g.ndata['api_idx'] = torch.tensor(apis_indices, dtype=torch.long, device=self.device) # Create node features

        thread_ids = sub_df['tid'].tolist()
        g.ndata['tid'] = torch.tensor(thread_ids, dtype=torch.long, device=self.device) # Add thread IDs as node features

        edges_src, edges_dst = [], []
        for tid in sub_df['tid'].unique():
            tid_mask = sub_df['tid'] == tid
            tid_indices = np.where(tid_mask)[0]
            if len(tid_indices) > 1:
                edges_src.extend(tid_indices[:-1])
                edges_dst.extend(tid_indices[1:])

        g.add_edges(
            torch.tensor(edges_src, dtype=torch.long, device=self.device),
            torch.tensor(edges_dst, dtype=torch.long, device=self.device)
        )

        label = sub_df['label'].iloc[0]
        g.ndata['label'] = torch.tensor([label] * g.num_nodes(), dtype=torch.long, device=self.device)
        return g, torch.tensor(label, dtype=torch.long, device=self.device)

```

```

def collate_samples(samples):
    graphs, labels = map(list, zip(*samples))
    batched_graph = dgl.batch(graphs)
    batched_labels = torch.stack(labels)
    return batched_graph, batched_labels

```

```

train_dataset = CustomGraphDataset(train_data, api_to_idx, device)

```

```
val_dataset = CustomGraphDataset(val_data, api_to_idx, device)
```

```
test_dataset = CustomGraphDataset(test_data, api_to_idx, device)
```

```
train_loader = GraphDataLoader(train_dataset, batch_size=16, shuffle=True, collate_fn=collate_samples)
```

```
val_loader = GraphDataLoader(val_dataset, batch_size=16, shuffle=True, collate_fn=collate_samples)
```

```
test_loader = GraphDataLoader(test_dataset, batch_size=16, shuffle=True, collate_fn=collate_samples)
```

GCN Model

```
class GraphModel(nn.Module):
    def __init__(self, num_apis, embedding_dim, num_classes, allow_zero_in_degree):
        super(GraphModel, self).__init__()
        self.embedding = nn.Embedding(num_apis, embedding_dim)
        self.conv1 = GraphConv(embedding_dim, 16, allow_zero_in_degree=allow_zero_in_degree)
        self.conv2 = GraphConv(16, num_classes, allow_zero_in_degree=allow_zero_in_degree)
        self.classifier = nn.Linear(num_classes, num_classes)# This layer is used to further process the aggregated graph-level features

    def forward(self, g, return_embeds=False):
        x = self.embedding(g.ndata['api_idx'])
        x = F.relu(self.conv1(g, x))
        x = self.conv2(g, x)
        g.ndata['h'] = x # Store the characteristics of each node
        hg = dgl.mean_nodes(g, 'h')# Average the features across all nodes to get graph level features
        if return_embeds:
            return hg # Return aggregated features for each graph
        return self.classifier(hg)# Classification of graph-level features
```

Train

```

def train(model, train_loader, val_loader, optimizer, criterion, device, num_epochs, patience):
    model.to(device)
    best_val_loss = np.inf
    epochs_no_improve = 0 # 用来跟踪验证损失是否已经停止改善

    for epoch in range(num_epochs):
        model.train() # 设置模型为训练模式
        total_loss = 0

        for batched_graph, labels in train_loader:
            batched_graph = batched_graph.to(device)
            labels = labels.to(device)

            optimizer.zero_grad()
            logits = model(batched_graph)
            loss = criterion(logits, labels)
            loss.backward()
            optimizer.step()
            total_loss += loss.item()

        average_train_loss = total_loss / len(train_loader)

        # 验证步骤
        model.eval() # 设置模型为评估模式
        total_val_loss = 0
        with torch.no_grad():
            for batched_graph, labels in val_loader:
                batched_graph = batched_graph.to(device)
                labels = labels.to(device)
                logits = model(batched_graph)
                val_loss = criterion(logits, labels)
                total_val_loss += val_loss.item()

        average_val_loss = total_val_loss / len(val_loader)
        print(f'Epoch {epoch + 1}/{num_epochs}, Train Loss: {average_train_loss:.4f}, Val Loss: {average_val_loss:.4f}')

        # 检查早停条件
        if average_val_loss < best_val_loss:
            best_val_loss = average_val_loss
            epochs_no_improve = 0
            torch.save(model.state_dict(), 'best_model.pth') # 保存最好的模型
        else:
            epochs_no_improve += 1
            if epochs_no_improve == patience:
                print(f'Early stopping triggered after {epoch + 1} epochs!')
                break # 停止训练

    # 加载最佳模型
    model.load_state_dict(torch.load('best_model.pth'))

```

```

num_apis = len(api_to_idx) # API索引的数量
model = GraphModel(num_apis=num_apis, embedding_dim=10, num_classes=8, allow_zero_in_degree=True)

```

```
optimizer = optim.Adam(model.parameters(), lr=0.001)
criterion = nn.CrossEntropyLoss()
num_epochs = 20
patience = 10
```

```
train(model, train_loader, val_loader, optimizer, criterion, device, num_epochs, patience)
```

```
model_save_path = '/content/drive/MyDrive/Project/best_model.pth'
```

```
torch.save(model.state_dict(), model_save_path)
```

```
model.load_state_dict(torch.load(model_save_path))
```



<All keys matched successfully>

Test

```
def evaluate(dataloader):
    model.eval()
    total_correct = 0
    total = 0
    with torch.no_grad():
        for batched_graph, labels in dataloader:
            batched_graph = batched_graph.to(device)
            labels = labels.to(device)

            outputs = model(batched_graph)
            _, predicted = torch.max(outputs, 1)
            total_correct += (predicted == labels).sum().item()
            total += labels.size(0)

    return total_correct / total
```

```
model.to(device)
test_accuracy = evaluate(test_loader)
```

```
print(f"Test Accuracy: {test_accuracy:.4f}")
```



Test Accuracy: 0.7558

```
def extract_features(dataloader):
    model.eval()
    all_features = []
    all_labels = []
    with torch.no_grad():
        for batched_graph, labels in dataloader:
            batched_graph = batched_graph.to(device)
            features = model(batched_graph, return_embeds=True) # 提取特征
            all_features.append(features.cpu().numpy()) # 存储特征
            all_labels.append(labels.cpu().numpy()) # 存储标签

    return np.concatenate(all_features), np.concatenate(all_labels)
```

```
train_features, train_labels = extract_features(train_loader)
val_features, val_labels = extract_features(val_loader)
test_features, test_labels = extract_features(test_loader)
```

```
def save_features_label(features_path, labels_path, features, labels):
    features_path = '/content/drive/MyDrive/Project/features/' + features_path
    labels_path = '/content/drive/MyDrive/Project/labels/' + labels_path

    np.save(features_path, features)
    np.save(labels_path, labels)
```

```
save_features_label('train_features.npy', 'train_labels.npy', train_features, train_labels)
```


```
save_features_label('val_features.npy', 'val_labels.npy', val_features, val_labels)
save_features_label('test_features.npy', 'test_labels.npy', test_features, test_labels)
```

```
def load_features_label(features_path, labels_path):
    features_path = '/content/drive/MyDrive/Project/features/' + features_path
    labels_path = '/content/drive/MyDrive/Project/labels/' + labels_path


    features = np.load(features_path)
    labels = np.load(labels_path)
    return features, labels
```

```
train_features, train_labels = load_features_label('train_features.npy', 'train_labels.npy')
val_features, val_labels = load_features_label('val_features.npy', 'val_labels.npy')
test_features, test_labels = load_features_label('test_features.npy', 'test_labels.npy')
```

train_labels

 array([0, 5, 1, ..., 0, 4, 0])

np.unique(train_labels)

 array([0, 1, 2, 3, 4, 5, 6, 7])

train_features

```
↵ array([[ -2.7863426 ,  1.0929561 ,  0.19866112, ...,  1.5189173 ,
          1.3572495 ,  2.8466122 ],
        [ -1.6078061 ,  0.34045684, -0.3316099 , ...,  1.9464468 ,
          0.0520192 ,  0.7229677 ],
        [ -0.38655454, -0.93857294, -0.48164195, ...,  0.68538797,
          -1.1652526 , -0.7453965 ],
        ...,
        [ -1.4786471 ,  0.5126234 ,  0.45795298, ...,  1.8214818 ,
          0.8521397 ,  2.5416162 ],
        [ -0.04052667, -2.005582 ,  0.38119894, ..., -0.06919952,
          -1.6395158 ,  0.09374472],
        [ -2.7388089 ,  1.1816062 ,  0.2791612 , ...,  2.2478292 ,
          1.0634265 ,  2.4594162 ]], dtype=float32)
```

```
def train_classifier(model, parameter_dict, x_train, y_train):
    # Use GridSearchCV to find the best hyperparameters
    grid_search = GridSearchCV(model, parameter_dict, cv=5, scoring=make_scorer(accuracy_score))
    grid_search.fit(x_train, y_train)

    # Print the best parameter and best score obtained
    print(f'The Best Parameter: {grid_search.best_params_}')
    print(f'The Best Score: {grid_search.best_score_}')

    return grid_search.best_estimator_
```

```
def test_confusion_matrix(model, x_test, y_test, train_labels):
    y_pred = model.predict(x_test)

    cm = confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(8, 6))
    plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
    plt.title('Confusion Matrix')
    plt.colorbar()
    tick_marks = np.arange(len(np.unique(train_labels)))
    plt.xticks(tick_marks, np.unique(train_labels), rotation=45)
    plt.yticks(tick_marks, np.unique(train_labels))
    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.show()
```

```
def save_model(model, filename):

    dump(model, filename)
    print(f"Model saved to {filename}")
```

```
def load_model(filename):

    model = load(filename)
    print(f"Model loaded from {filename}")
    return model
```

regression, decision Tree, Boost, GLM, CTree, Random Forest, Artificial Neural Network

LogisticRegression

```
scaler = StandardScaler()
train_features_scaled = scaler.fit_transform(train_features)
log_params = {
    'fit_intercept': [True, False],
    'C': [0.1, 1, 10],
    'solver': ['lbfgs', 'saga'],
    'max_iter': [1000] # 增加迭代次数
}
best_clf = train_classifier(LogisticRegression(), log_params, train_features_scaled, train_labels)
#1m8s
```

→ /usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
The Best Parameter: {'C': 10, 'fit_intercept': True, 'max_iter': 1000, 'solver': 'saga'}
The Best Score: 0.7551440329218108
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(

```
test_features_scaled = scaler.fit_transform(test_features)
y_pred_lr = best_clf.predict(test_features_scaled)
```


```
print(classification_report(test_labels, y_pred_lr, target_names=np.unique(train_labels).astype(str)))
```

→


	precision	recall	f1-score	support
0	0.83	0.87	0.85	745
1	0.86	0.80	0.83	74
2	0.79	0.78	0.79	190
3	0.48	0.69	0.57	116
4	0.00	0.00	0.00	14
5	0.80	0.80	0.80	635
6	0.45	0.16	0.24	79
7	0.52	0.48	0.50	231
accuracy			0.75	2084
macro avg	0.59	0.57	0.57	2084
weighted avg	0.74	0.75	0.74	2084

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted sample
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted sample
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted sample
_warn_prf(average, modifier, msg_start, len(result))

test_confusion_matrix(best_clf, test_features_scaled, test_labels, train_labels)



```
test_ROC_curve_multiclass(best_clf, test_features_scaled, test_labels, np.unique(train_labels))
```



RandomForestClassifier

```
rf_params = {'n_estimators': [100, 200], 'max_depth': [None, 10, 20, 30]}  
best_rf = train_classifier(RandomForestClassifier(random_state=42), rf_params, train_features, train_labels)
```

The Best Parameter: {'max_depth': 20, 'n_estimators': 200}
The Best Score: 0.8452674897119342

```
y_pred_rf = best_rf.predict(test_features)
```

```
accuracy = accuracy_score(test_labels, y_pred_rf)  
print("Accuracy:", accuracy)
```


Accuracy: 0.8498080614203455

```
print(classification_report(test_labels, y_pred_rf, target_names=np.unique(train_labels).astype(str)))
```


↔

	precision	recall	f1-score	support
0	0.89	0.95	0.92	745
1	0.93	0.93	0.93	74
2	0.92	0.88	0.90	190
3	0.56	0.77	0.64	116
4	0.67	0.14	0.24	14
5	0.92	0.89	0.90	635
6	0.61	0.52	0.56	79
7	0.68	0.55	0.61	231
accuracy			0.85	2084
macro avg	0.77	0.70	0.71	2084
weighted avg	0.85	0.85	0.85	2084

```
test_confusion_matrix(best_rf, test_features, test_labels, train_labels)
```



```
test_ROC_curve_multiclass(best_rf, test_features, test_labels, np.unique(train_labels))
```



DecisionTreeClassifier

```
dt_params = {  
    'max_depth': [None, 10, 20, 30],  
    'min_samples_split': [2, 10, 20],  
    'min_samples_leaf': [1, 5, 10]  
}  
best_dt = train_classifier(DecisionTreeClassifier(random_state=42), dt_params, train_features, train_labels)
```

The Best Parameter: {'max_depth': 10, 'min_samples_leaf': 1, 'min_samples_split': 2}
The Best Score: 0.7988683127572016


```
y_pred_dt = best_dt.predict(test_features)
```

```
print(classification_report(test_labels, y_pred_dt, target_names=np.unique(train_labels).astype(str)))
```

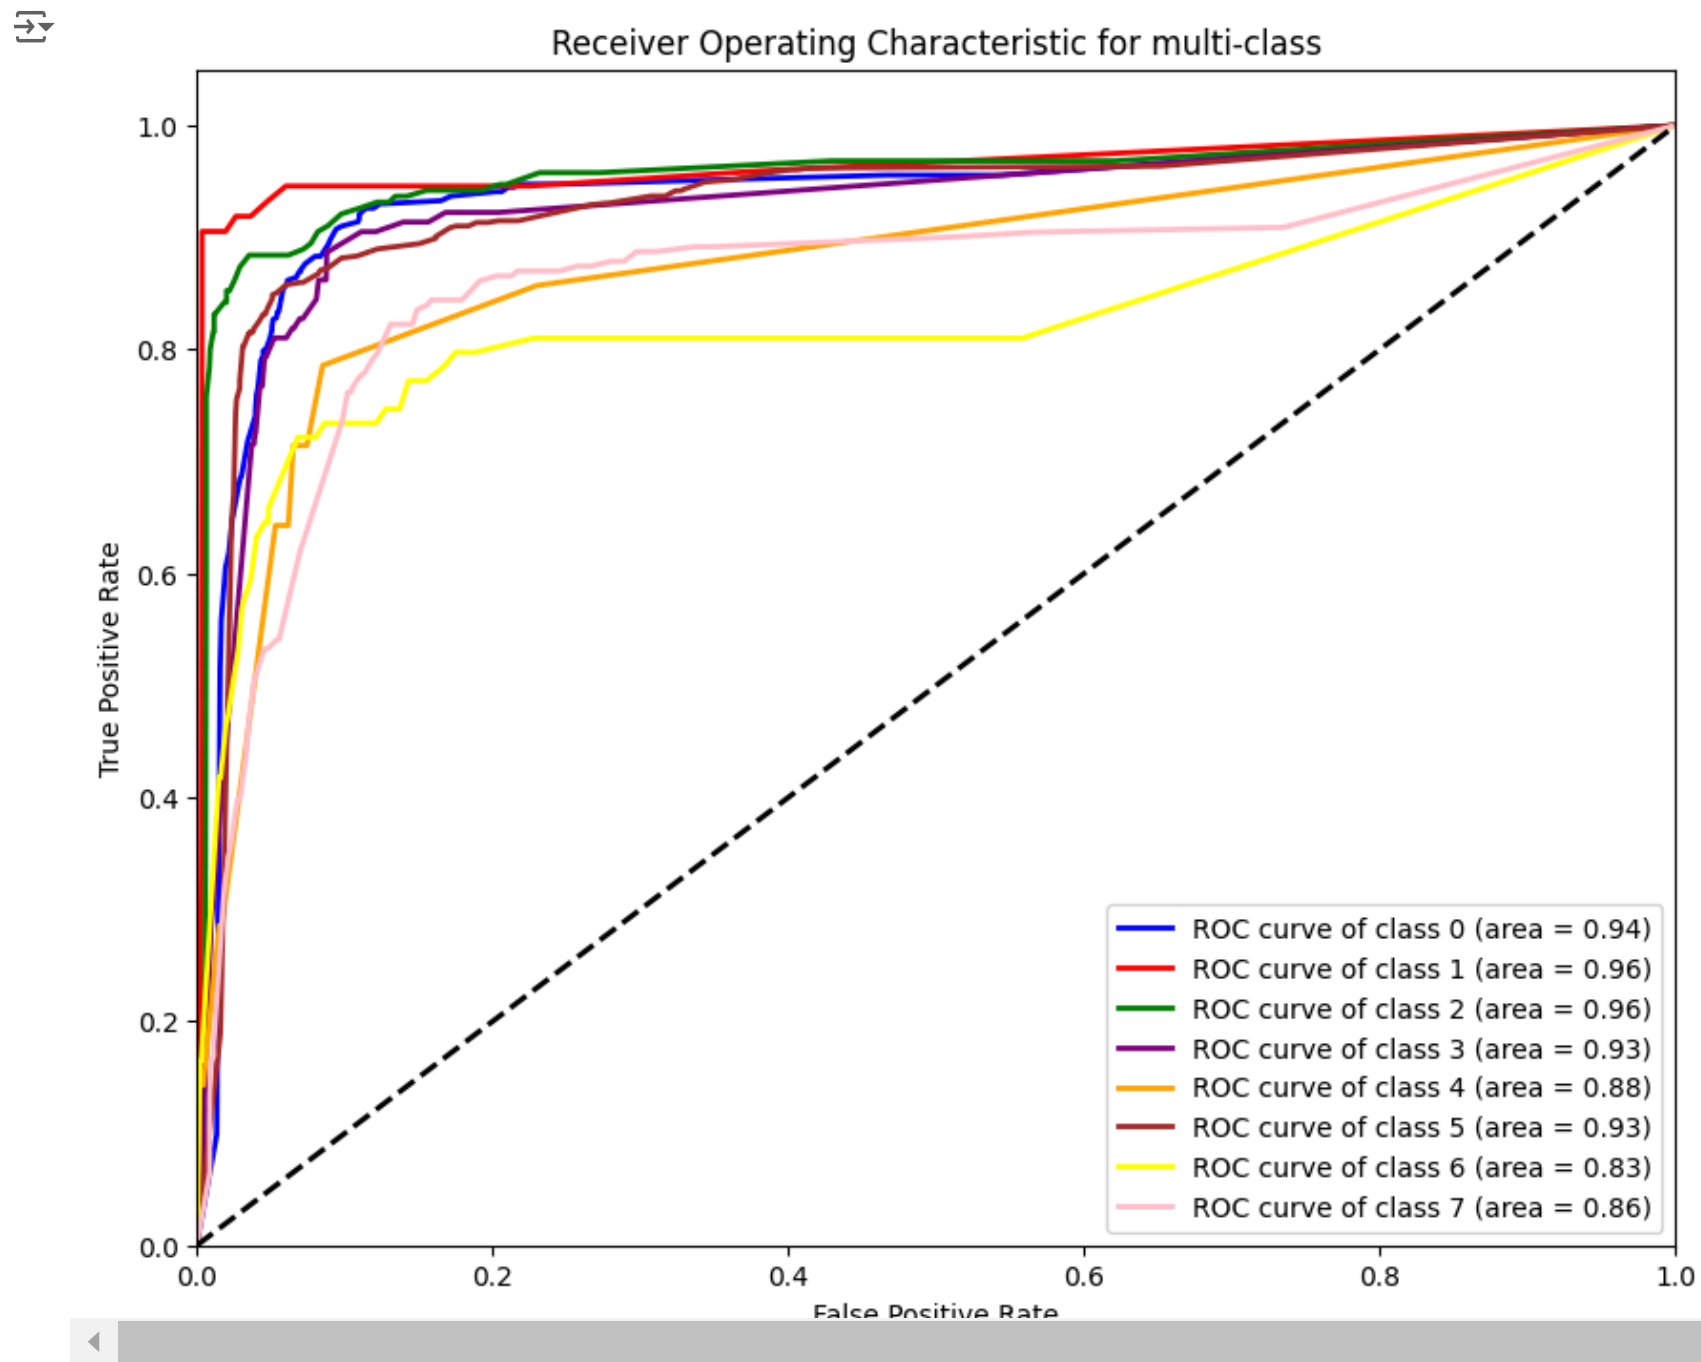


	precision	recall	f1-score	support
0	0.85	0.89	0.87	745
1	0.90	0.88	0.89	74
2	0.80	0.85	0.83	190
3	0.49	0.73	0.59	116
4	0.40	0.14	0.21	14
5	0.90	0.82	0.86	635
6	0.52	0.43	0.47	79
7	0.59	0.52	0.55	231
accuracy			0.79	2084
macro avg	0.68	0.66	0.66	2084
weighted avg	0.80	0.79	0.79	2084

```
test_confusion_matrix(best_dt, test_features, test_labels, train_labels)
```



```
test_ROC_curve_multiclass(best_dt, test_features, test_labels, np.unique(train_labels))
```




GradientBoostingClassifier

```
gb_params = {  
    'n_estimators': [100, 200], # Number of trees  
    'learning_rate': [0.05, 0.1], # learning rate  
    'max_depth': [3, 5], # Maximum depth  
    'min_samples_split': [2], # Minimum Sample Segmentation  
    'min_samples_leaf': [1] # Minimum sample leaf node  
}  
# Setting up early stops  
best_GBC = train_classifier(GradientBoostingClassifier(random_state=42, n_estimators=300, validation_fraction=0.1, n_iter_no_change=10, tol=0.01), gb_params, train_features, train_labels)
```

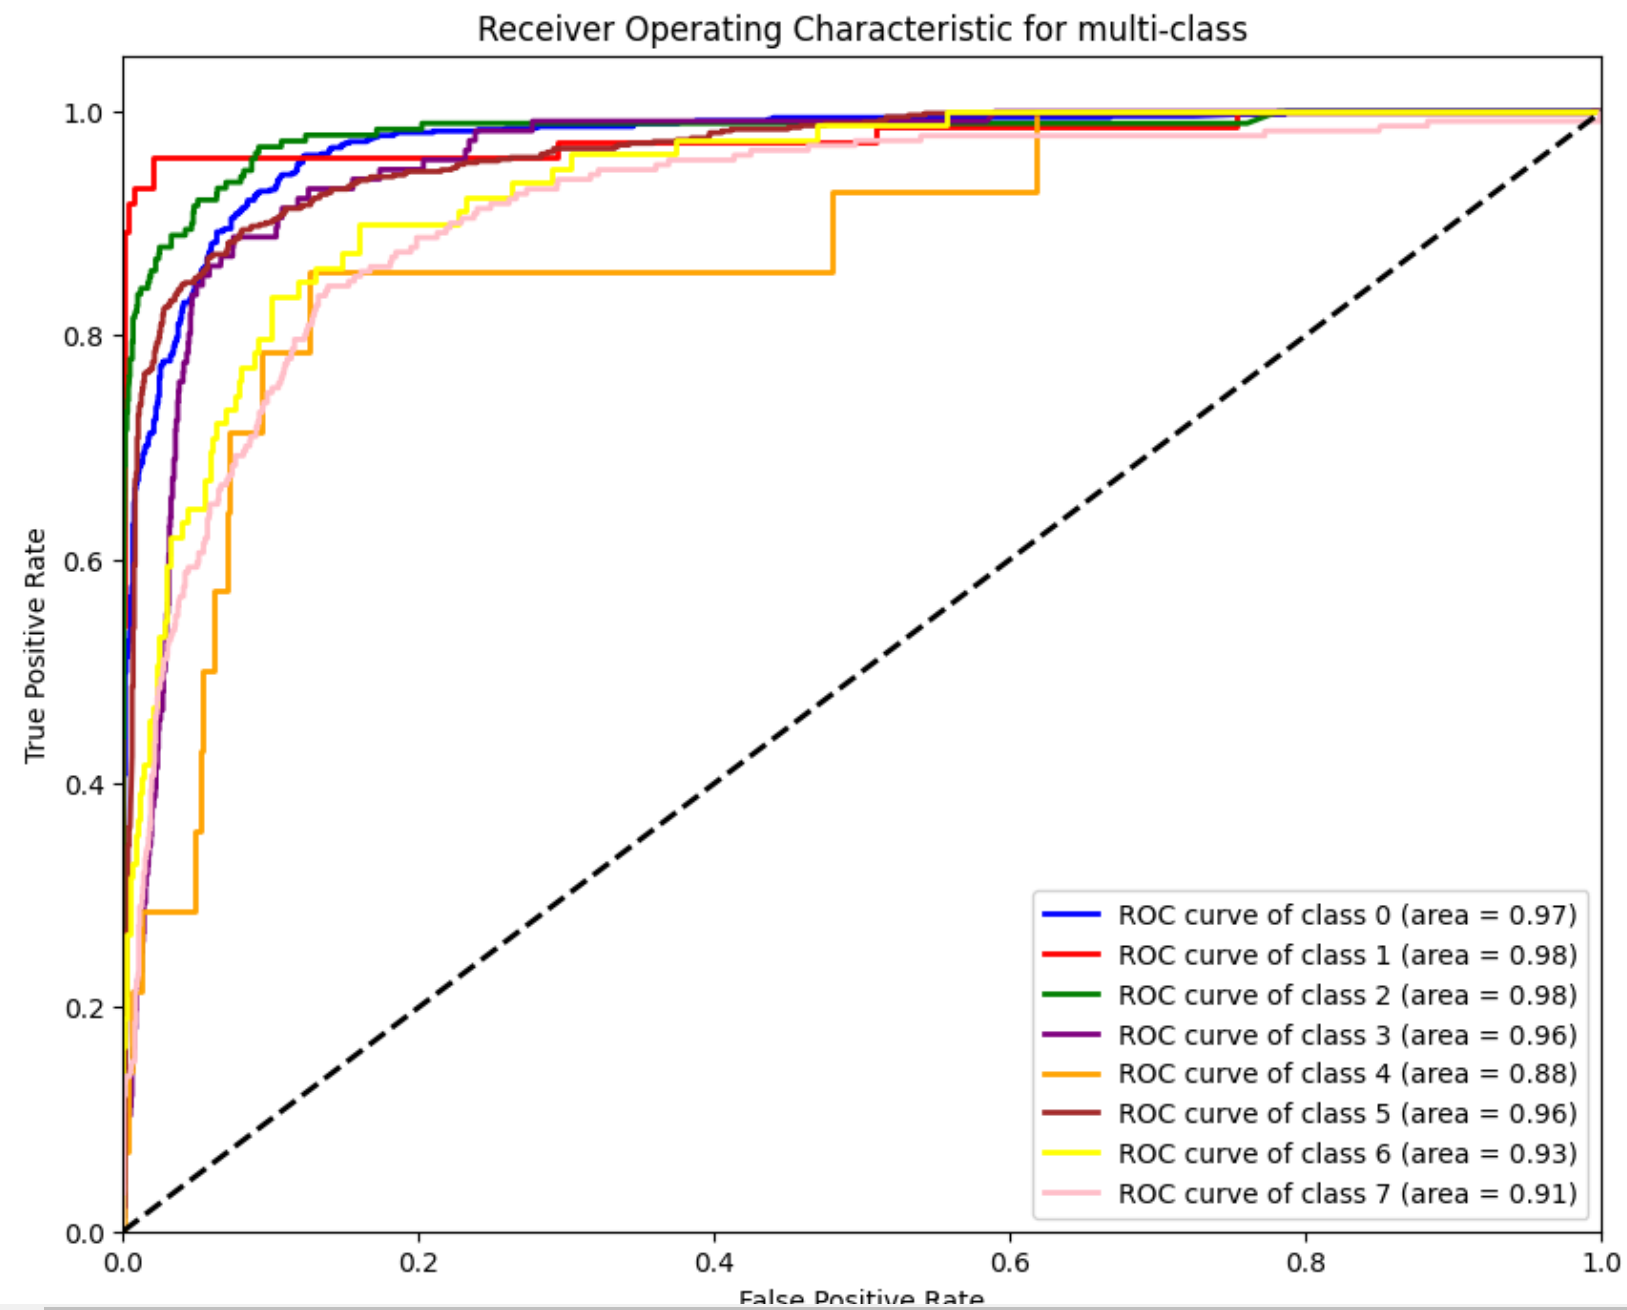
The Best Parameter: {'learning_rate': 0.1, 'max_depth': 5, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 100}
The Best Score: 0.8083333333333333

```
y_pred_gbc = best_GBC.predict(test_features)
```

```
test_confusion_matrix(best_GBC, test_features, test_labels, train_labels)
```



```
test_ROC_curve_multiclass(best_GBC, test_features, test_labels, np.unique(train_labels))
```

MLP

```
mlp_params = {
    'hidden_layer_sizes': [(50,), (100,), (50, 50)],
    'activation': ['tanh', 'relu'],
    'solver': ['sgd', 'adam'],
    'learning_rate_init': [0.001, 0.01],
    'max_iter': [500, 1000]
}
best_ann = train_classifier(MLPClassifier(random_state=42), mlp_params, train_features, train_labels)
```




4	0.00	0.00	0.00	14
5	0.90	0.86	0.88	635
6	0.52	0.47	0.49	79
7	0.63	0.54	0.58	231
accuracy			0.82	2084
macro avg	0.65	0.66	0.65	2084
weighted avg	0.82	0.82	0.82	2084

```

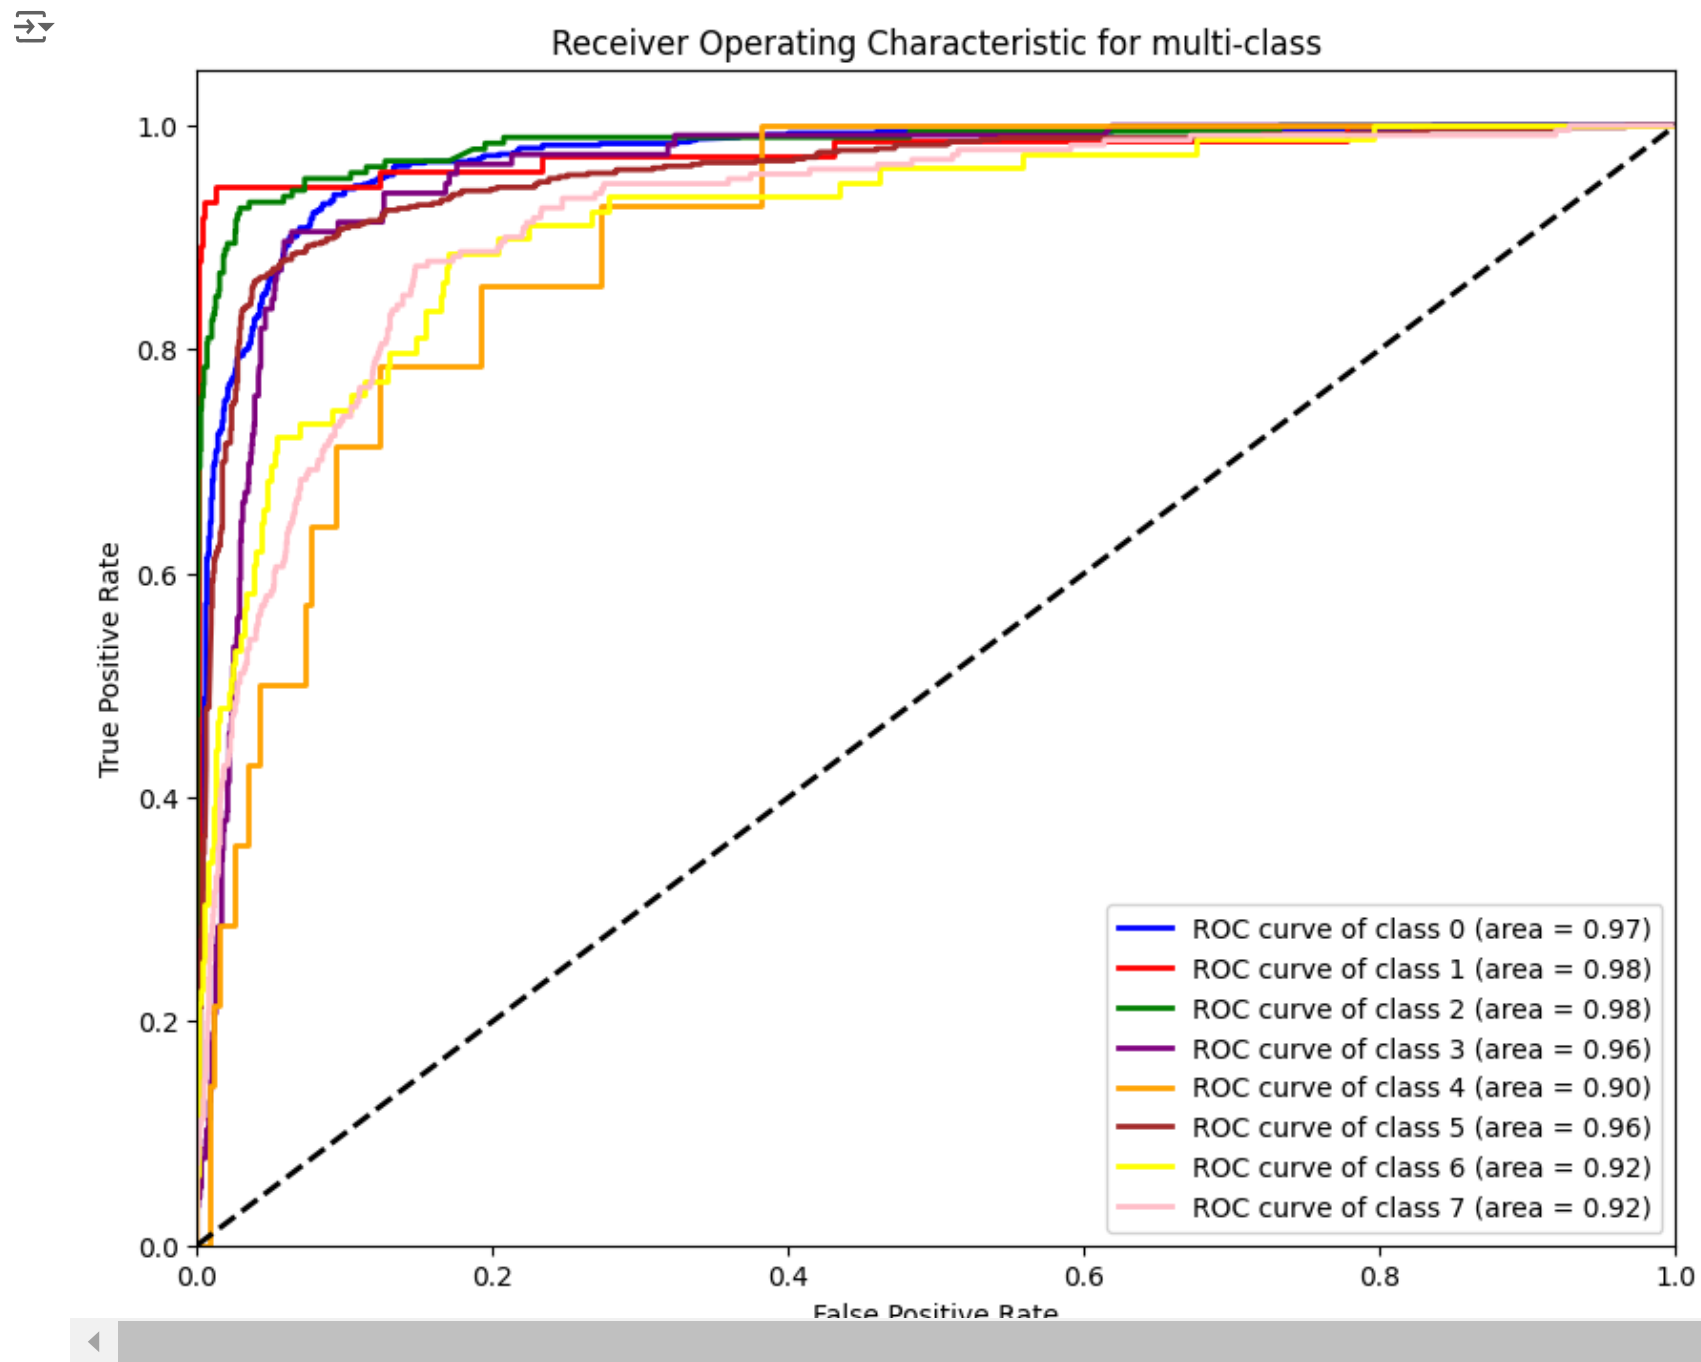
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted sample
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted sample
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted sample
_warn_prf(average, modifier, msg_start, len(result))

```

```
test_confusion_matrix(best_ann, test_features, test_labels, train_labels)
```



```
test_ROC_curve_multiclass(best_ann, test_features, test_labels, np.unique(train_labels))
```



Evaluation

```
report_lr = classification_report(test_labels, y_pred_lr, target_names=np.unique(train_labels).astype(str), output_dict=True)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted sample
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted sample
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted sample
_warn_prf(average, modifier, msg_start, len(result))
```

```
report_rf = classification_report(test_labels, y_pred_rf, target_names=np.unique(train_labels).astype(str), output_dict=True)
```

```
report_dt = classification_report(test_labels, y_pred_dt, target_names=np.unique(train_labels).astype(str), output_dict=True)
```

```
report_ann = classification_report(test_labels, y_pred_ann, target_names=np.unique(train_labels).astype(str), output_dict=True)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted sample
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted sample
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted sample
_warn_prf(average, modifier, msg_start, len(result))
```

```
report_gbc = classification_report(test_labels, y_pred_gbc , target_names=np.unique(train_labels).astype(str), output_dict=True)
```

```
df_lr = pd.DataFrame(report_lr).transpose()
df_rf = pd.DataFrame(report_rf).transpose()
df_dt = pd.DataFrame(report_dt).transpose()
df_ann = pd.DataFrame(report_ann).transpose()
```

```
df_gbc = pd.DataFrame(report_gbc).transpose()
```

```
dfs = [df_lr, df_rf, df_dt, df_ann, df_gbc]
for df in dfs:
    df.set_index(df.index.str.strip(), inplace=True)
df_combined = pd.concat(dfs, axis=1, keys=['LR', 'RF', 'DT', 'ANN', 'GBC'])

print(df_combined)
```

	LR				RF		\
	precision	recall	f1-score	support	precision	recall	
0	0.828863	0.871141	0.849476	745.000000	0.894073	0.951678	
1	0.855072	0.797297	0.825175	74.000000	0.932432	0.932432	
2	0.792553	0.784211	0.788360	190.000000	0.918033	0.884211	
3	0.484848	0.689655	0.569395	116.000000	0.556250	0.767241	
4	0.000000	0.000000	0.000000	14.000000	0.666667	0.142857	
5	0.795597	0.796850	0.796223	635.000000	0.917342	0.891339	
6	0.448276	0.164557	0.240741	79.000000	0.611940	0.518987	
7	0.518692	0.480519	0.498876	231.000000	0.679144	0.549784	
accuracy	0.751919	0.751919	0.751919	0.751919	0.849808	0.849808	
macro avg	0.590488	0.573029	0.571031	2084.000000	0.771985	0.704816	
weighted avg	0.742822	0.751919	0.743581	2084.000000	0.849859	0.849808	

	DT				\
	f1-score	support	precision	recall	f1-score
0	0.921977	745.000000	0.849490	0.893960	0.871158
1	0.932432	74.000000	0.902778	0.878378	0.890411
2	0.900804	190.000000	0.801980	0.852632	0.826531
3	0.644928	116.000000	0.494186	0.732759	0.590278
4	0.235294	14.000000	0.400000	0.142857	0.210526
5	0.904153	635.000000	0.898451	0.822047	0.858553
6	0.561644	79.000000	0.515152	0.430380	0.468966
7	0.607656	231.000000	0.594059	0.519481	0.554273
accuracy	0.849808	0.849808	0.794626	0.794626	0.794626
macro avg	0.713611	2084.000000	0.682012	0.659062	0.658837
weighted avg	0.846452	2084.000000	0.798185	0.794626	0.793488

	ANN				\
	support	precision	recall	f1-score	support
0	745.000000	0.879896	0.904698	0.892124	745.000000
1	74.000000	0.929577	0.891892	0.910345	74.000000

2	190.000000	0.830846	0.878947	0.854220	190.000000
3	116.000000	0.517857	0.750000	0.612676	116.000000
4	14.000000	0.000000	0.000000	0.000000	14.000000
5	635.000000	0.901478	0.864567	0.882637	635.000000
6	79.000000	0.521127	0.468354	0.493333	79.000000
7	231.000000	0.631313	0.541126	0.582751	231.000000
accuracy	0.794626	0.818138	0.818138	0.818138	0.818138
macro avg	2084.000000	0.651512	0.662448	0.653511	2084.000000
weighted avg	2084.000000	0.816547	0.818138	0.815467	2084.000000

	GBC				
	precision	recall	f1-score	support	
0	0.859671	0.912752	0.885417	745.0000	
1	0.883117	0.918919	0.900662	74.0000	
2	0.883978	0.842105	0.862534	190.0000	
3	0.542857	0.655172	0.593750	116.0000	
4	0.166667	0.142857	0.153846	14.0000	
5	0.876404	0.859843	0.868045	635.0000	
6	0.525424	0.392405	0.449275	79.0000	
7	0.641791	0.558442	0.597222	231.0000	
accuracy	0.811900	0.811900	0.811900	0.8119	
macro avg	0.672489	0.660312	0.663844	2084.0000	
weighted avg	0.808707	0.811900	0.808951	2084.0000	

```
def combined_graphs(df_combined, str):
    parameter_lr = df_combined['LR'][str]
    parameter_rf = df_combined['RF'][str]
    parameter_dt = df_combined['DT'][str]
    parameter_ann = df_combined['ANN'][str]
    parameter_gbc = df_combined['GBC'][str]

    df_parameters = pd.DataFrame({
        'LR': parameter_lr,
        'RF': parameter_rf,
        'DT': parameter_dt,
        'ANN': parameter_ann,
        'GBC': parameter_gbc
    })

    ax = df_parameters.plot(kind='bar', figsize=(14, 8))
    ax.set_xlabel('Class')
    ax.set_ylabel(str)
    ax.set_title(str + 'Comparison Across Different Models')
    ax.legend(title='Models')


    plt.xticks(rotation=0)
    plt.tight_layout()
    plt.show()

combined_graphs(df_combined, 'precision')
```


```
combined_graphs(df_combined, 'precision')
```



precisionComparison Across Different Models




```
combined_graphs(df_combined, 'recall')
```





recallComparison Across Different Models



```
combined_graphs(df_combined, 'f1-score')
```


f1-scoreComparison Across Different Models

